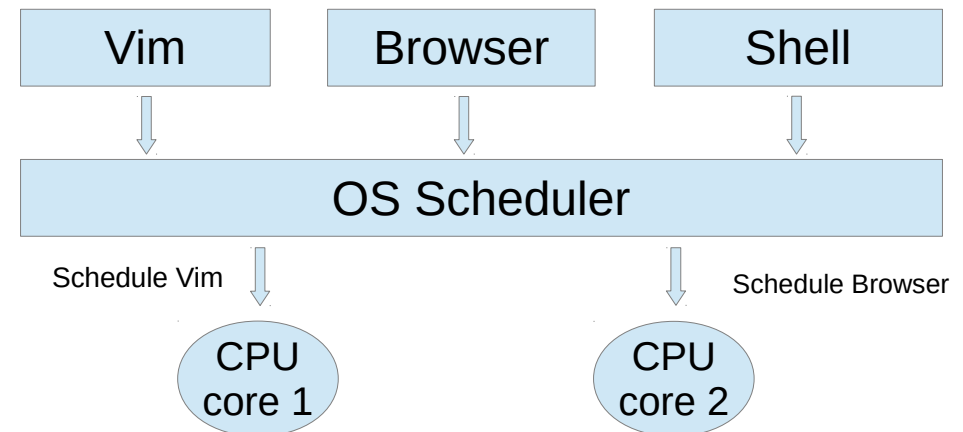
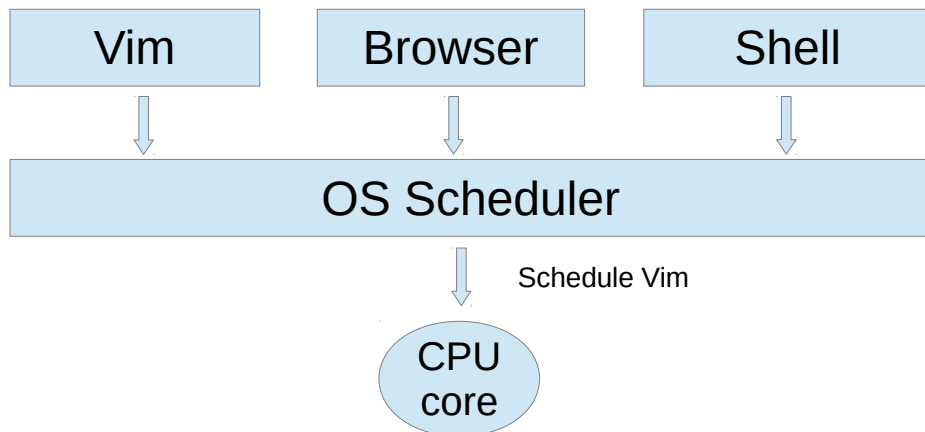


# Multithreading/Parallel Processing

Week 8

# Multitasking

- Run multiple processes **simultaneously** to increase performance
- Processes do not share internal structures (stack, globals, etc)
  - Communicate via **IPC** (inter-process communication) methods
    - Pipes, Sockets, Signals, Message Queues, etc
- **Single-core: Illusion** of parallelism by switching processes quickly (**time-sharing**)
- **Multi-core: True** parallelism. Multiple processes execute **concurrently** on different CPU cores



# Multitasking

- `tr -s '[:space:]' '\n' | sort -u | comm -23 - words`
- Three separate processes spawned **simultaneously**
  - P1 – `tr`
  - P2 – `sort`
  - P3 – `comm`
- Common buffers (**pipes**) exist between 2 processes for communication
  - '`tr`' writes its `stdout` to a buffer that is read by '`sort`'
  - '`sort`' can execute, as and when data is available in the buffer
  - Similarly, a buffer is used for communicating between '`sort`' and '`comm`'

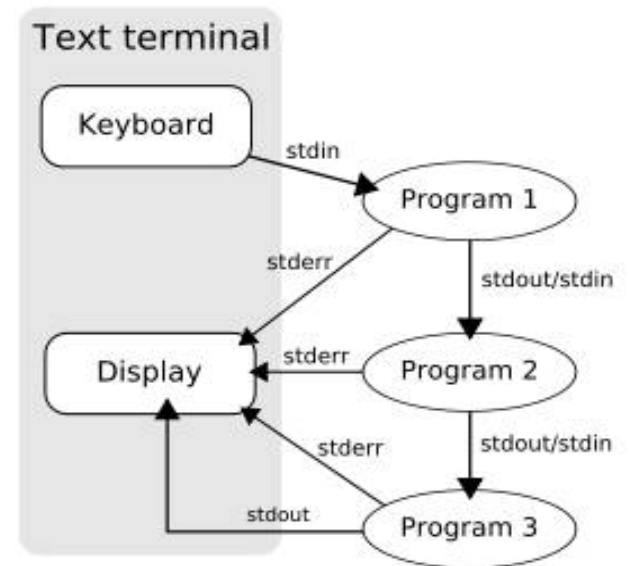
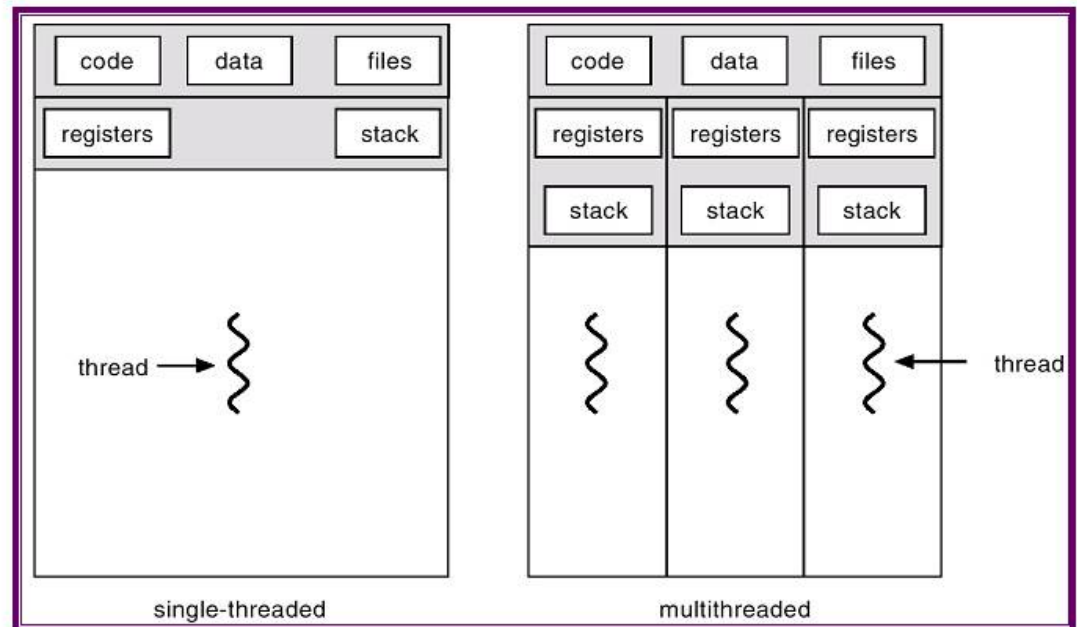


Image source: Wikipedia

# Threads

- A process can be
  - **Single-threaded**
  - **Multi-threaded**
- Threads in a process can run in **parallel**
- A thread is a **lightweight process**
- It is a **basic unit** of CPU utilization
- Each thread has its **own**
  - Stack
  - Registers
  - Thread ID
- Each thread **shares** the following with other threads belonging to the same process
  - Code
  - Global Data
  - OS resources (files, I/O)



# Single threaded execution

- `int globalCounter = 0;`

```
int main()
```

```
{
```

```
...
```

```
foo(arg1, arg2);
```

```
bar(arg3, arg4, arg5);
```

```
...
```

```
return 0;
```

```
}
```

- `void foo(arg1, arg2)`

```
{
```

```
//code for foo
```

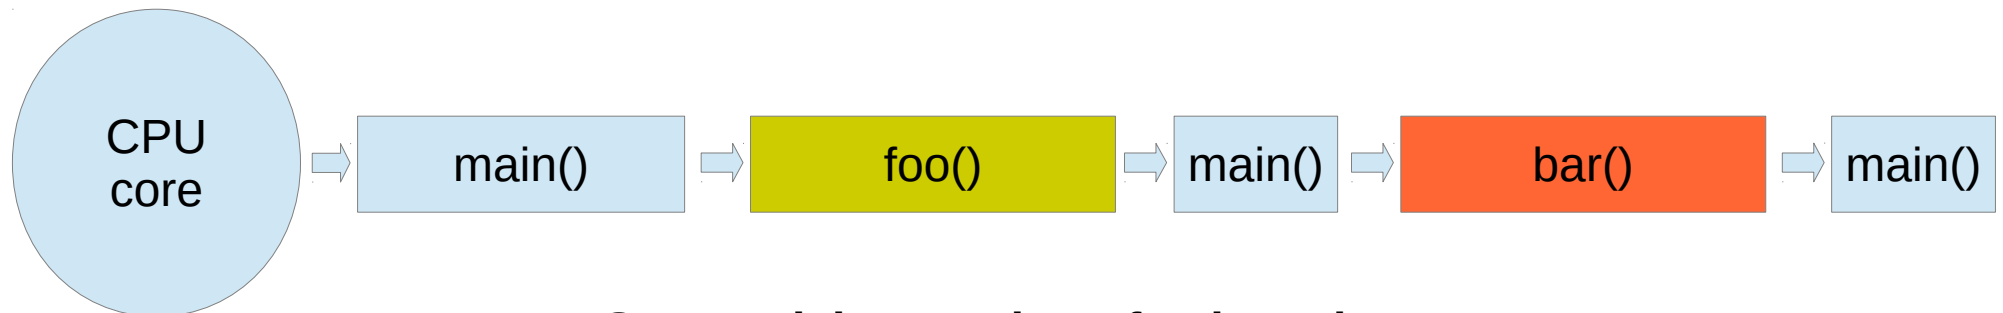
```
}
```

- `void bar(arg3, arg4, arg5)`

```
{
```

```
//code for bar
```

```
}
```



**Sequential execution of subroutines**

# Multi-threaded execution (single core)

- `int globalCounter = 0;`

```
int main()
```

```
{
```

```
...
```

```
Run thread foo(arg1, arg2);
```

```
Run thread bar(arg3, arg4, arg5);
```

```
...
```

```
return 0;
```

```
}
```

- `void foo(arg1, arg2)`

```
{
```

```
//code for foo
```

```
}
```

- `void bar(arg3, arg4, arg5)`

```
{
```

```
//code for bar
```

```
}
```



**Time Sharing – Illusion of multithreaded parallelism**  
(Thread switching has less overhead compared to process switching)

# Multi-threaded execution

- `int globalCounter = 0;`

`int main()`

`{`

`...`

`Run thread foo(arg1, arg2);`

`Run thread bar(arg3, arg4, arg5);`

`...`

`return 0;`

`}`

- `void foo(arg1, arg2)`

`{`

`//code for foo`


`}`

- `void bar(arg3, arg4, arg5)`

`{`

`//code for bar`

`}`



**True multithreaded parallelism**

# Multithreading properties

- Efficient way to **parallelize** tasks
- **Thread switches are less expensive** compared to process switches (context switching)
- Inter-thread communication is easy, via **shared global** data
- Need **synchronization** among threads accessing same data



# Pthread API

```
#include <pthread.h>
```

- `int pthread_create(pthread_t *thread,  
                    const pthread_attr_t *attr,  
                    void* (*thread_function)(void *),  
                    void *arg);`
  - Returns 0 on success, otherwise returns non-zero error number
- `void pthread_exit(void *retval);`
- `int pthread_join(pthread_t thread, void **retval);`
  - Returns 0 on success, otherwise returns non-zero error number

# Pthreads sample code

Compile the following code as `gcc main.c -lpthread`

```
#include<pthread.h>
#include<stdio.h>

void* ThreadFunction(void *arg)
{
    long tID = (long)arg;
    printf("Inside thread function with ID = %ld\n", tID);
    pthread_exit(0);
}

int main(int argc, char *argv[])
{
    const int nthreads = 5;
    pthread_t threadID[nthreads];
    long t;
    for(t = 0; t < nthreads; ++t) {
        int rs = pthread_create(&threadID[t], 0, ThreadFunction, (void*)t);
        if(rs) {
            fprintf(stderr, "Error creating thread\n");
            return -1;
        }
    }

    printf("Main thread finished creating threads\n");

    for(t = 0; t < nthreads; ++t) {
        void *retVal;
        int rs = pthread_join(threadID[t], &retVal);
        if(rs) {
            fprintf(stderr, "Error joining thread\n");
            return -1;
        }
    }

    printf("Main thread finished execution!\n");

    return 0;
}
```

# Thread safety/synchronization

- **Thread safe function** – Safe to be called by multiple threads at the same time. Function is free of 'race conditions' when called by multiple threads simultaneously.
- **Race condition** – The output depends on the order of execution
  - Shared data changed by 2 threads
    - `int balance = 1000;`
  - Thread 1
    - T1 - Read balance
    - T1 - Deduct 50 from balance
    - T1 - Update balance with new value
  - Thread 2
    - T2 - Read balance
    - T2 - Add 150 to balance
    - T2 - Update balance with new value

# Thread synchronization

- **Order 1**

- balance = 1000
- T1 – Read balance (1000)
- T1 – Deduct 50
  - 950 in temporary result
- T2 – Read balance (1000)
- T1 – Update balance
  - balance is 950 at this point
- T2 – Add 150 to balance
  - 1150 in temporary result
- T2 – Update balance
  - balance is 1150 at this point

- **The final value of balance is 1150**

- **Order 2**

- balance = 1000
- T1 – Read balance (1000)
- T2 – Read balance (1000)
- T2 – Add 150 to balance
  - 1150 in temporary result
- T1 – Deduct 50
  - 950 in temporary result
- T2 – Update balance
  - balance is 1150 at this point
- T1 – Update balance
  - balance is 950 at this point

- **The final value of balance is 950**

# Thread Synchronization

- **Mutex (Mutual exclusion)**

- Thread 1

- Mutex.lock()

- Read balance

- Deduct 50 from balance

- Update balance with new value

- Mutex.unlock()

- Thread 2

- Mutex.lock()

- Read balance

- Add 150 to balance

- Update balance with new value

- Mutex.unlock()

- balance = 1100

- Only one thread will get the mutex. Other threads will **block in Mutex.lock()**.
- Other threads can start execution only when the thread having the mutex calls **Mutex.unlock()**

# Lab

- Evaluate performance of multithreaded 'sort' command
- `od -An -f -N 4000000 < /dev/urandom | tr -s ' ' '\n' > random.txt`
- You might have to modify the above command
- Delete empty line
- `time -p sort -g --parallel=2 numbers.txt > /dev/null`

# Ray-Tracing

- **Powerful rendering technique in Computer Graphics**
- **Yields very high quality rendering**
  - Suited for scenes with complex light interactions
  - Visually realistic
  - Trace the path of light in the scene
- **Computationally very expensive**
  - Not suited for rendering in **real-time** (example:games)
  - Suited for rendering high-quality pictures
- **Embarrassingly parallel**
  - Good candidate for **multi-threading**
  - Threads need **not synchronize** with each other, because each thread works on a different pixel

# Ray-tracing






Image Source: POV Ray, Hall of Fame



# Ray-tracing

- Trace the path of a ray from the eye
  - **One ray per pixel** in the view window
  - The color of the ray is the color of the corresponding pixel
- Check for **intersection** of ray with scene objects.
- **Lighting**
  - **Flat shading** – The whole object has uniform brightness
  - **Lambertian shading** – Cosine of angle between surface normal and light direction



# Homework Output



# Homework - Anti-Aliasing

